

## dense\_retriever

April 27, 2025

```
[1]: # %%
import os
import json
import pickle
import numpy as np
from tqdm import tqdm
import torch
from transformers import AutoTokenizer, AutoModel
import faiss
from huggingface_hub import login
from config import config
login(config['HF_API_KEY'])

class DenseRetriever:
    """
    Dense retriever implementation using a pre-trained language model
    for encoding documents and queries into vector representations.
    """

    def __init__(self, model_name="nlpueb/legal-bert-base-uncased",
        ↪index_name="legal_dense_index"):
        """
        Initialize the dense retriever with a pre-trained language model.

        Args:
            model_name: Name of the pre-trained model to use for encoding
            index_name: Name for the index directory
        """
        self.model_name = model_name
        self.index_name = index_name
        self.index_dir = os.path.join(os.getcwd(), index_name)
        os.makedirs(self.index_dir, exist_ok=True)

        # Initialize model and tokenizer
        self.tokenizer = None
        self.model = None
        self.index = None
```

```

self.documents = None
self.doc_ids = None
self.embedding_dim = None

def _initialize_model(self):
    """Load the model and tokenizer if not already loaded"""
    if self.tokenizer is None:
        print(f>Loading tokenizer for {self.model_name}...)
        self.tokenizer = AutoTokenizer.from_pretrained(self.model_name)

    if self.model is None:
        print(f>Loading model {self.model_name}...)
        self.model = AutoModel.from_pretrained(self.model_name)
        self.model.eval() # Set to evaluation mode

def get_embedding(self, text, max_length=512):
    """Generate embedding for a single text"""
    # Ensure model and tokenizer are loaded
    self._initialize_model()

    # Tokenize input
    inputs = self.tokenizer(text, return_tensors="pt",
↪max_length=max_length,
                                padding="max_length", truncation=True)

    # Generate embeddings
    with torch.no_grad():
        outputs = self.model(**inputs)
        # Use CLS token embedding as text representation
        embedding = outputs.last_hidden_state[:, 0, :].cpu().numpy()

    return embedding[0] # Return as 1D array

def index_corpus(self, documents, doc_ids, batch_size=8):
    """
    Generate embeddings for all documents and build a search index.

    Args:
    documents: List of document texts
    doc_ids: List of document IDs corresponding to the documents
    batch_size: Batch size for processing documents
    """
    print(f>Building dense index with {len(documents)} documents...)

    # Store document texts and IDs
    self.documents = documents
    self.doc_ids = doc_ids

```

```

# Ensure model and tokenizer are loaded
self._initialize_model()

# Generate embeddings for all documents
print("Generating document embeddings...")
all_embeddings = []

for i in tqdm(range(0, len(documents), batch_size), desc="Processing_
↳document batches"):
    batch_docs = documents[i:i+batch_size]
    batch_inputs = self.tokenizer(batch_docs, padding=True,
↳truncation=True,
                                return_tensors="pt", max_length=512)

    with torch.no_grad():
        outputs = self.model(**batch_inputs)
        # Use CLS token embedding
        batch_embeddings = outputs.last_hidden_state[:, 0, :].cpu().
↳numpy()
        all_embeddings.append(batch_embeddings)

# Concatenate all batch embeddings
document_embeddings = np.vstack(all_embeddings)
self.embedding_dim = document_embeddings.shape[1]

# Normalize embeddings for cosine similarity
faiss.normalize_L2(document_embeddings)

# Build FAISS index for fast similarity search
print(f"Building FAISS index with {document_embeddings.shape[1]}_
↳dimensions...")
self.index = faiss.IndexFlatIP(document_embeddings.shape[1]) # Inner_
↳product for cosine similarity
self.index.add(document_embeddings)

# Save the index and metadata
self.save_index()

print("Dense index built successfully")
return self

def save_index(self):
    """Save the index and associated data to disk"""
    print(f"Saving index to {self.index_dir}...")

    # Save FAISS index

```

```

        if self.index is not None:
            faiss.write_index(self.index, os.path.join(self.index_dir, "faiss.
↪index"))

        # Save metadata
        metadata = {
            "model_name": self.model_name,
            "embedding_dim": self.embedding_dim,
            "num_documents": len(self.documents) if self.documents else 0
        }

        with open(os.path.join(self.index_dir, "metadata.json"), 'w') as f:
            json.dump(metadata, f, indent=2)

        # Save documents and doc_ids
        with open(os.path.join(self.index_dir, "documents.json"), 'w') as f:
            json.dump(self.documents, f)

        with open(os.path.join(self.index_dir, "doc_ids.json"), 'w') as f:
            json.dump(self.doc_ids, f)

    def load_index(self):
        """Load pre-built index and associated data"""
        index_path = os.path.join(self.index_dir, "faiss.index")
        if not os.path.exists(index_path):
            raise ValueError(f"Index not found at {index_path}. Build index_
↪first with index_corpus()")

        print(f"Loading index from {self.index_dir}...")

        # Load FAISS index
        self.index = faiss.read_index(index_path)

        # Load metadata
        with open(os.path.join(self.index_dir, "metadata.json"), 'r') as f:
            metadata = json.load(f)
            self.model_name = metadata["model_name"]
            self.embedding_dim = metadata["embedding_dim"]

        # Load documents and doc_ids
        with open(os.path.join(self.index_dir, "documents.json"), 'r') as f:
            self.documents = json.load(f)

        with open(os.path.join(self.index_dir, "doc_ids.json"), 'r') as f:
            self.doc_ids = json.load(f)

        print(f"Loaded dense index with {len(self.documents)} documents")

```

```

    return self

def retrieve(self, query, k=100):
    """
    Retrieve top-k documents for a query.

    Args:
        query: Query text
        k: Number of documents to retrieve

    Returns:
        List of dictionaries with document ID, score, and text
    """
    if self.index is None:
        self.load_index()

    # Generate query embedding
    query_embedding = self.get_embedding(query)
    query_embedding = query_embedding.reshape(1, -1)

    # Normalize query embedding for cosine similarity
    faiss.normalize_L2(query_embedding)

    # Search index
    scores, indices = self.index.search(query_embedding, k)

    # Format results
    results = []
    for i, idx in enumerate(indices[0]):
        if idx < len(self.documents): # Safety check
            results.append({
                "id": self.doc_ids[idx],
                "score": float(scores[0][i]),
                "text": self.documents[idx]
            })

    return results

def batch_retrieve(self, queries, k=100):
    """
    Retrieve top-k documents for multiple queries.

    Args:
        queries: List of query texts
        k: Number of documents to retrieve per query

    Returns:

```

```

        Dictionary mapping query index to list of results
    """
    if self.index is None:
        self.load_index()

    all_results = {}
    for i, query in enumerate(tqdm(queries, desc="Processing queries")):
        all_results[str(i)] = self.retrieve(query, k=k)

    return all_results

def save_results(self, results, output_file):
    """Save retrieval results to file"""
    with open(output_file, 'w') as f:
        json.dump(results, f, indent=2)
    print(f"Saved retrieval results to {output_file}")

# %%
# Usage example with your generated corpus
if __name__ == "__main__":
    # Load the generated corpus
    corpus_file = "legal_dummy_corpus.json"
    queries_file = "legal_sample_queries.json"

    # Check if corpus file exists
    if not os.path.exists(corpus_file):
        print("Corpus file not found. Please generate the corpus first.")
        exit(1)
    else:
        # Load existing corpus
        print(f"Loading corpus from {corpus_file}...")
        with open(corpus_file, 'r') as f:
            corpus_data = json.load(f)
            documents = corpus_data["documents"]
            doc_ids = corpus_data["doc_ids"]

    # Check if queries file exists
    if not os.path.exists(queries_file):
        print("Queries file not found. Please generate the queries first.")
        exit(1)
    else:
        # Load existing queries
        print(f"Loading queries from {queries_file}...")
        with open(queries_file, 'r') as f:
            queries = json.load(f)

    print(f"Corpus has {len(documents)} documents")

```

```

print(f"Query set has {len(queries)} queries")

# Initialize Dense retriever with a legal-domain model
# Using nlpaueb/legal-bert-base-uncased - specialized for legal text
retriever = DenseRetriever(
    model_name="nlpaueb/legal-bert-base-uncased",
    index_name="legal_dense_retr2"
)

# Check if index already exists
if os.path.exists(os.path.join(retriever.index_dir, "faiss.index")):
    print("Dense index already exists. Loading...")
    retriever.load_index()
else:
    print("Building new dense index...")
    retriever.index_corpus(documents, doc_ids)

# Retrieve results for queries
results = retriever.batch_retrieve(queries, k=10)

# Save results
output_file = "dense_retrieval_results.json"
retriever.save_results(results, output_file)

# Print sample results
print("\nSample Retrieval Results:")
print("=====")

for i, query in enumerate(queries[:3]): # Show results for first 3 queries
    print(f"\nQuery: {query}")
    print("-" * 80)
    results_for_query = results[str(i)]

    for j, doc in enumerate(results_for_query[:2]): # Show top 2 documents
        print(f"Document {j+1}: (Score: {doc['score']:.4f})")
        print(f"ID: {doc['id']}")
        print(f"Text: {doc['text'][:200]}..." if len(doc['text']) > 200
else f"Text: {doc['text']}")
        print("-" * 40)

```

```

/home/bandham/miniconda3/envs/llm692_venv/lib/python3.13/site-
packages/tqdm/auto.py:21: TqdmWarning: IProgress not found. Please update
jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm

```

```

Loading corpus from legal_dummy_corpus.json...
Loading queries from legal_sample_queries.json...

```

Corpus has 120 documents  
Query set has 15 queries  
Building new dense index...  
Building dense index with 120 documents...  
Loading tokenizer for nlpaueb/legal-bert-base-uncased...  
Loading model nlpaueb/legal-bert-base-uncased...  
Generating document embeddings...  
  
Processing document batches: 100%| | 15/15 [00:39<00:00, 2.61s/it]  
  
Building FAISS index with 768 dimensions...  
Saving index to /home/bandham/Documents/StonyBrook\_CourseWork/Spring 2025/LLM-AMS692.02/Legal-Mind/legal\_dense\_retr2...  
Dense index built successfully  
  
Processing queries: 100%| | 15/15 [00:36<00:00, 2.46s/it]  
  
Saved retrieval results to dense\_retrieval\_results.json  
  
Sample Retrieval Results:  
=====

Query: What are the essential elements of a valid contract?  
-----

Document 1: (Score: 0.9432)  
ID: criminal\_law\_030  
Text: SEARCH WARRANT application states that probable cause exists to believe evidence of possession with intent to distribute will be found at Northern University Campus based on DNA analysis observed by O...  
-----

Document 2: (Score: 0.9333)  
ID: administrative\_law\_018  
Text: AGENCY DECISION: Securities and Exchange Commission hereby approves/denies {party}'s application for operating license based on findings that the party demonstrated financial responsibility. This deci...  
-----

Query: How is negligence defined in tort law?  
-----

Document 1: (Score: 0.5427)  
ID: contract\_law\_028  
Text: EMPLOYMENT CONTRACT: Smith Corp. agrees to employ Commonwealth of Jefferson as Chief Financial Officer commencing on September 9, 2022 for three years. Compensation shall be \$300,000 per annum with be...  
-----

Document 2: (Score: 0.4883)  
ID: property\_law\_092  
Text: EASEMENT: MediCorp grants to PacificRoute Services a perpetual easement for conservation over the property described as a 20-foot wide strip along the western edge of the property. This easement shall...



-----  
Query: What constitutes probable cause for a search warrant?  
-----

Document 1: (Score: 0.5912)

ID: contract\_law\_028

Text: EMPLOYMENT CONTRACT: Smith Corp. agrees to employ Commonwealth of Jefferson as Chief Financial Officer commencing on September 9, 2022 for three years. Compensation shall be \$300,000 per annum with be...

-----  
Document 2: (Score: 0.5856)

ID: criminal\_law\_013

Text: INDICTMENT: The Grand Jury charges that on July 8, 2022, defendant Richard Taylor did knowingly and intentionally misrepresented material facts, constituting the offense of money laundering under §9.2...

## 0.1 USAGE

```
[3]: # from dense_retriever import DenseRetriever # Import your class

retriever = DenseRetriever(
    model_name="nlpaueb/legal-bert-base-uncased",
    index_name="legal_dense_retr2"
)

retriever.load_index()

# 3. Search with any new query
results = retriever.retrieve("What are the essential elements of a valid_
↪contract?", k=5)

# 4. Process the results
for i, result in enumerate(results):
    print(f"Result {i+1}: {result['text'][:100]}... (Score: {result['score']:.
↪4f})")
```

Loading index from /home/bandham/Documents/StonyBrook\_CourseWork/Spring  
2025/LLM-AMS692.02/Legal-Mind/legal\_dense\_retr2...

Loaded dense index with 120 documents

Loading tokenizer for nlpaueb/legal-bert-base-uncased...

Loading model nlpaueb/legal-bert-base-uncased...

Result 1: SEARCH WARRANT application states that probable cause exists to  
believe evidence of possession with ... (Score: 0.9432)

Result 2: AGENCY DECISION: Securities and Exchange Commission hereby  
approves/denies {party}'s application for... (Score: 0.9333)

Result 3: PLEA AGREEMENT: Defendant Jennifer Lee, charged with criminal

negligence, agrees to plead guilty to ... (Score: 0.9310)

Result 4: In In re Wilson Estate (2022), the Court held that religious freedom protected under the Sixth Amend... (Score: 0.9307)

Result 5: CONSTITUTIONAL ANALYSIS: The Senate Bill 247 must be subjected to heightened scrutiny under the Four... (Score: 0.9304)